

- When conducting the work of Lab 11, we conducted the test that uses the Central Limit Theorem even though the sample size was “small” (i.e., $n < 30$). It turns out, that how “far off” the t -test is can be computed using a first-order Edgeworth approximation for the error. Below, we will do this for the the further observations.

(a) Boos and Hughes-Oliver (2000) note that

$$P(T \leq t) \approx F_Z(t) + \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6} f_Z(t)}_{\text{error}},$$

where $f_Z(\cdot)$ and $F_Z(\cdot)$ are the Gaussian PDF and CDF and skew is the skewness of the data. What is the potential error in the computation of the p -value when testing $H_0 : \mu_X = 0; H_a : \mu_X < 0$ using the zebra finch further data?

```
#####
# PART A
#####
#initializing data
finch_data = read_csv("zebrafinches.csv")

## Rows: 25 Columns: 3
## -- Column specification -----
## Delimiter: ","
## dbl (3): closer, further, diff
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.

further_data = finch_data$`further`
closer_data = finch_data$`closer`
diff_data = finch_data$`diff`
skew_further = skewness(further_data)
n=nrow(finch_data)
#Calculating t statistic
t_val = mean(further_data)/(sd(further_data)/sqrt(n))

#Calculating potential error
potential_err = (skew_further/sqrt(n))*((2*t_val^2+1)/6)*dnorm(t_val)
```

- Compute the error for t statistics from -10 to 10 and plot a line that shows the error across t . Continue to use the skewness and the sample size for the zebra finch further data.

```
#####
# PART B
#####
#Initializing
t_vals = seq(from = -10, to = 10, by = 0.01)
R = length(t_vals)
err_tibble = tibble(err=rep(NA,R))

#Looping over each t value
for (i in 1:R){
  t_val = t_vals[i]
  #Calculating potential error
  potential_err = (skew_further/sqrt(n))*((2*t_val^2+1)/6)*dnorm(t_val)

  #Adding error to tibble
  err_tibble$err[i] = potential_err
}
#view(err_tibble)

#Making plot of values
error_plot = ggplot(data=err_tibble, aes(t_vals, err)) +
  geom_line() +
  theme_bw() +
  geom_hline(yintercept = 0) +
  xlab("T Value") +
  ylab("Potential Error")
```




Figure 1: Error Across t

- (c) Suppose we wanted to have a tail probability within 10% of the desired $\alpha = 0.05$. Recall we did a left-tailed test using the further data. How large of a sample size would we need? That is, we need to solve the error formula equal to 10% of the desired left-tail probability:

$$0.10\alpha \stackrel{\text{set}}{=} \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6} f_Z(t)}_{\text{error}},$$

which yields

$$n = \left(\frac{\text{skew}}{6(0.10\alpha)} (2t^2 + 1) f_Z(t) \right)^2.$$

```
#####
# PART C
#####
alpha = 0.05
t_val = qnorm(alpha)
n = ((skew_further/(6*0.1*alpha)) * (2*t_val^2+1)*dnorm(t_val))^2
```

2. Complete the following steps to revisit the analyses from lab 11 using the bootstrap procedure.

- (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform resampling to approximate the sampling distribution of the T statistic:

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}},$$

where \bar{x}_r is the mean computed on the r^{th} resample and s is the sample standard deviation from the original samples. At the end, create an object called `resamples.null.closer`, for example, and store the resamples shifted to ensure they are consistent with the null hypotheses at the average (i.e., here ensure the shifted resamples are 0 on average, corresponding to $t = 0$, for each case).

```
R <- 10000
resamples <- tibble(further = rep(NA, R),
                    closer = rep(NA, R),
                    diff = rep(NA, R))

for(i in 1:R){
  #Making a sample for further
```

```

curr.resample <- sample(further_data,
                        size = length(further_data),
                        replace = T)

resamples$furthur[i] <- mean(curr.resample)/(sd(further_data)/sqrt(length(further_data)))

#Making a sample for closer
curr.resample <- sample(closer_data,
                        size = length(closer_data),
                        replace = T)

resamples$closer[i] <- mean(curr.resample)/(sd(closer_data)/sqrt(length(closer_data)))

#Making a sample for difference
curr.resample <- sample(diff_data,
                        size = length(diff_data),
                        replace = T)

resamples$diff[i] <- mean(curr.resample)/(sd(diff_data)/sqrt(length(diff_data)))
}
resamples_shifted = resamples |>
  mutate(further = further - mean(further_data)/(sd(further_data)/sqrt(25))) |>
  mutate(closer = closer - mean(closer_data)/(sd(closer_data)/sqrt(25))) |>
  mutate(diff = diff - mean(diff_data)/(sd(diff_data)/sqrt(25)))

#view(resamples)
#view(resamples_shifted)

```

- (b) Compute the bootstrap p -value for each test using the shifted resamples. How do these compare to the t -test p -values?

```

#####
# PART B
#####
pval_further = (mean(resamples_shifted$furthur <= mean(resamples$furthur)))/R
pval_closer = (mean(resamples_shifted$closer >= mean(resamples$diff)))/R
pval_diff = (mean(resamples_shifted$diff >= mean(resamples$diff)))/R

pval_further

## [1] 0

pval_closer

## [1] 0

pval_diff

## [1] 0

```

- (c) What is the 5th percentile of the shifted resamples under the null hypothesis? Note this value approximates $t_{0.05, n-1}$. Compare these values in each case.

```

#####
# PART C
#####
resample_further_5th = quantile(resamples_shifted$furthur, c(0.05, 0.95))
resample_closer_5th = quantile(resamples_shifted$closer, c(0.05, 0.95))
resample_diff_5th = quantile(resamples_shifted$diff, c(0.05, 0.95))

resample_further_5th

##          5%          95%
## -1.693523  1.588579

resample_closer_5th

##          5%          95%
## -1.598645  1.668597

resample_diff_5th

##          5%          95%
## -1.539859  1.668092

```

- (d) Compute the bootstrap confidence intervals using the resamples. How do these compare to the t -test confidence intervals?

```
# Confidence Interval
resample_further_ci = quantile(resamples$further, c(0.025, 0.975))
resample_closer_ci = quantile(resamples$closer, c(0.025, 0.975))
resample_diff_ci = quantile(resamples$diff, c(0.025, 0.975))

resample_further_ci

##      2.5%      97.5%
## -9.817407 -5.964390

resample_closer_ci

##      2.5%      97.5%
##  6.413167 10.280816

resample_diff_ci

##      2.5%      97.5%
##  6.683322 10.520937
```

3. Complete the following steps to revisit the analyses from lab 11 using the randomization procedure.

- (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform the randomization procedure

```
#####
# PART A
#####
closer_data = finch_data$closer
diff_data = finch_data$diff

R <- 10000
rand <- tibble(further = rep(NA, R),
               closer = rep(NA, R),
               diff = rep(NA, R))

#Randomization for further
# PREPROCESSING: shift the data to be mean 0 under H0
x.shift <- further_data - 0
# RANDOMIZE / SHUFFLE
for(i in 1:R){
  curr.rand <- x.shift *
    sample(x = c(-1, 1),
           size = length(x.shift),
           replace = T)

  rand$further[i] <- mean(curr.rand)
}

#Randomization for closer
x.shift <- closer_data - 0
# RANDOMIZE / SHUFFLE
for(i in 1:R){
  curr.rand <- x.shift *
    sample(x = c(-1, 1),
           size = length(x.shift),
           replace = T)

  rand$closer[i] <- mean(curr.rand)
}

#Randomization for difference
x.shift <- diff_data - 0
# RANDOMIZE / SHUFFLE
for(i in 1:R){
  curr.rand <- x.shift *
    sample(x = c(-1, 1),
           size = length(x.shift),
           replace = T)

  rand$diff[i] <- mean(curr.rand)
}
#view(rand)
```

- (b) Compute the randomization test p -value for each test.

```
#####
# PART B
#####
# p-value for further
(delta <- abs(mean(further_data) - 0))

## [1] 0.2027244

(low <- 0 - delta) # mirror

## [1] -0.2027244

(high<- 0 + delta) # xbar

## [1] 0.2027244

pval_further = mean(rand$further <= low) + mean(rand$further >= high)

# p-value for closer
(delta <- abs(mean(closer_data) - 0))

## [1] 0.1562231

(low <- 0 - delta) # mirror

## [1] -0.1562231

(high<- 0 + delta) # xbar

## [1] 0.1562231

pval_closer = mean(rand$closer <= low) + mean(rand$closer >= high)

# p-value for further
(delta <- abs(mean(diff_data) - 0))

## [1] 0.3589475

(low <- 0 - delta) # mirror

## [1] -0.3589475

(high<- 0 + delta) # xbar

## [1] 0.3589475

pval_diff = mean(rand$diff <= low) + mean(rand$diff >= high)
```

- (c) Compute the randomization confidence interval by iterating over values of μ_0 .

Hint: You can “search” for the lower bound from Q_1 and subtracting by 0.0001, and the upper bound using Q_3 and increasing by 0.0001. You will continue until you find the first value for which the two-sided p -value is greater than or equal to 0.05.

```
#####
# PART C
#####
R <- 1000
mu0.iterate <- 0.001
starting.point <- mean(further_data)

#Calculating lower value
mu.lower <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- further_data - mu.lower
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
             size = length(x.shift),
             replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
}
```

```

# Thinking is hard
rand <- rand |>
  mutate(xbars = xbars + mu.lower) # shifting back

# p-value
(delta <- abs(mean(further_data) - mu.lower))
(low <- mu.lower - delta) # mirror
(high <- mu.lower + delta) # xbar
(p.val <- mean(rand$xbars <= low) +
  mean(rand$xbars >= high))

if(p.val < 0.05){
  break
}else{
  mu.lower <- mu.lower - mu0.iterate
}
}

#Calculating upper value
mu.upper <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- further_data - mu.upper
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
        size = length(x.shift),
        replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.upper) # shifting back

  # p-value
  (delta <- abs(mean(further_data) - mu.upper))
  (low <- mu.upper - delta) # mirror
  (high <- mu.upper + delta) # xbar
  (p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high))

  if(p.val < 0.05){
    break
  }else{
    mu.upper <- mu.upper + mu0.iterate
  }
}

further_ci = c(mu.lower, mu.upper)

#####
# CLOSER
#####
starting.point <- mean(closer_data)

#Calculating lower value
mu.lower <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- closer_data - mu.lower
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
        size = length(x.shift),
        replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.lower) # shifting back

```

```

# p-value
(delta <- abs(mean(closer_data) - mu.lower))
(low <- mu.lower - delta) # mirror
(high <- mu.lower + delta) # xbar
(p.val <- mean(rand$xbars <= low) +
  mean(rand$xbars >= high))

if(p.val < 0.05){
  break
}else{
  mu.lower <- mu.lower - mu0.iterate
}
}

#Calculating upper value
mu.upper <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- closer_data - mu.upper
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
        size = length(x.shift),
        replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.upper) # shifting back

  # p-value
  (delta <- abs(mean(closer_data) - mu.upper))
  (low <- mu.upper - delta) # mirror
  (high <- mu.upper + delta) # xbar
  (p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high))

  if(p.val < 0.05){
    break
  }else{
    mu.upper <- mu.upper + mu0.iterate
  }
}

closer_ci = c(mu.lower, mu.upper)

#####
# DIFFERENCE
#####
starting.point <- mean(diff_data)

#Calculating lower value
mu.lower <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- diff_data - mu.lower
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
        size = length(x.shift),
        replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.lower) # shifting back

  # p-value
  (delta <- abs(mean(diff_data) - mu.lower))
  (low <- mu.lower - delta) # mirror

```

```

(high<- mu.lower + delta) # xbar
(p.val <- mean(rand$xbars <= low) +
  mean(rand$xbars >= high))

if(p.val < 0.05){
  break
}else{
  mu.lower <- mu.lower - mu0.iterate
}
}

#Calculating upper value
mu.upper <- starting.point
repeat{
  rand <- tibble(xbars = rep(NA, R))

  # PREPROCESSING: shift the data to be mean 0 under H0
  x.shift <- diff_data - mu.upper
  # RANDOMIZE / SHUFFLE
  for(i in 1:R){
    curr.rand <- x.shift *
      sample(x = c(-1, 1),
        size = length(x.shift),
        replace = T)

    rand$xbars[i] <- mean(curr.rand)
  }
  # Thinking is hard
  rand <- rand |>
    mutate(xbars = xbars + mu.upper) # shifting back

  # p-value
  (delta <- abs(mean(diff_data) - mu.upper))
  (low <- mu.upper - delta) # mirror
  (high<- mu.upper + delta) # xbar
  (p.val <- mean(rand$xbars <= low) +
    mean(rand$xbars >= high))

  if(p.val < 0.05){
    break
  }else{
    mu.upper <- mu.upper + mu0.iterate
  }
}

diff_ci = c(mu.lower, mu.upper)

further_ci

## [1] -0.2587244 -0.1497244

closer_ci

## [1] 0.1182231 0.1972231

diff_ci

## [1] 0.2709475 0.4459475

```

4. **Optional Challenge:** In this lab, you performed resampling to approximate the sampling distribution of the T statistic using

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}}.$$

I'm curious whether it is better/worse/similar if we computed the statistics using the sample standard deviation of the resamples (s_r), instead of the original sample (s)

$$T = \frac{\bar{x}_r - 0}{s_r/\sqrt{n}}.$$

- Perform a simulation study to evaluate the Type I error for conducting this hypothesis test both ways.
- Using the same test case(s) as part (a), compute bootstrap confidence intervals and assess their coverage – how often do we ‘capture’ the parameter of interest?

References

Boos, D. D. and Hughes-Oliver, J. M. (2000). How large does n have to be for z and t intervals? *The American Statistician*, 54(2):121–128.